

Light-curve classification in massive variability surveys – I. Microlensing

Vasily Belokurov,¹ N. Wyn Evans^{1,2*} and Yann Le Du¹

¹*Theoretical Physics, Department of Physics, 1 Keble Road, Oxford OX1 3NP*

²*Institute of Astronomy, Madingley Rd, Cambridge CB3 0HA*

Accepted 2003 February 7. Received 2003 February 6; in original form 2002 November 6

ABSTRACT

This paper exploits neural networks to provide a fast and automatic way to classify light curves in massive photometric data sets. As an example, we provide a working neural network that can distinguish microlensing light curves from other forms of variability, such as eruptive, pulsating, cataclysmic and eclipsing variable stars. The network has five input neurons, a hidden layer of five neurons and one output neuron. The five input variables for the network are extracted by spectral analysis from the light-curve data points and are optimized for the identification of a single, symmetric, microlensing bump. The output of the network is the posterior probability of microlensing.

The committee of neural networks successfully passes tests on noisy data taken by the MACHO collaboration. When used to process ~ 5000 light curves on a typical tile towards the bulge, the network cleanly identifies the single microlensing event. When fed with a subsample of 36 light curves identified by the MACHO collaboration as microlensing, the network corroborates this verdict in the case of 27 events, but classifies the remaining nine events as other forms of variability. For some of these discrepant events, it looks as though there are secondary bumps or the bump is noisy or not properly contained. Neural networks naturally allow for the possibility of novelty detection; that is, new or unexpected phenomena which we may want to follow-up. The advantages of neural networks for microlensing rate calculations, as well as the future developments of massive variability surveys, are both briefly discussed.

Key words: gravitational lensing – methods: data analysis – techniques: photometric – surveys – stars: variables: other.

1 INTRODUCTION

Variability in the sky has been known for thousands of years, but our understanding of variable sources remains very incomplete. Some of the most interesting objects in the sky are transient. These include supernovae (SNe), microlensed stars, near-Earth or killer asteroids (which are transient because of their exceptionally large proper motions) optical flashes associated with gamma-ray bursts and stars undergoing short-lived but key stages of stellar evolution such as the helium-core flash, etc. All of these objects are rare. To hunt them down in a systematic way means that we must record images, process the data in real time (or nearly so), recognize the events from their light curves and archive them.

The earliest examples of massive transient astronomy searches are the microlensing surveys such as MACHO,¹ EROS² and OGLE.³ Typically, the surveys monitored $\sim 5 \times 10^6$ stars a few times every

night over several years in the directions of the Galactic Bulge and the Magellanic clouds, yielding $\sim 10^{10}$ photometric measurements. Out of the $\sim 10^5$ sources that were variable, the surveys tried to identify $\sim 10^2$ true microlensing events. The selection criteria typically involved the imposition of sets of cuts to ensure good light-curve coverage and a steady baseline flux, to require a single bump and thus eliminate common forms of stellar variability and to require a good statistical fit to the achromatic standard microlensing light curve and so on. Many of the cuts were developed through trial and error, and evolved as the experiments progressed (e.g. Alcock et al. 1997, 2000a). Unambiguous identification of microlensing events was sometimes not possible, and the collaborations sometimes reported their results in terms of two sets, one of high-quality events (any light curve that was undoubtedly microlensing) and one of possible events (any light curve with a unique peak and a flat baseline). Sometimes the cuts even eliminated interesting events – for example, the longest-ever microlensing event OGLE-1999-BUL-32 was originally missed as its baseline flux was not constant and so failed one of the imposed cuts (Mao et al. 2002).

Additionally, microlensing alert or early-warning systems (e.g. Udalski et al. 1994) work by reducing the number of candidates to

*E-mail: w.evans1@physics.oxford.ac.uk

¹<http://www.macho.anu.edu.au/>

²<http://eros.in2p3.fr/>

³<http://sirius.astrouw.edu.pl/~ogle/>

manageable amounts. The candidates for each night are individually examined for the onset of microlensing. Even for surveys as large as OGLE II, this worked well. However, still larger surveys are planned for the future and therefore it becomes important to automate the procedure and issue alerts without human intervention.

The microlensing experiments are of course not the only massive photometry searches being conducted by astronomers at the moment. There are also collaborations primarily looking for supernovae (e.g. The Supernovae Cosmology Project), optical flashes related to gamma-ray bursts (*ROTSE*) and near-Earth asteroids (NEAT and LINEAR). More generally, as Paczyński (2001, 2002) has emphasized, the monitoring of the optical sky for variability is likely to enjoy a huge resurgence over the coming decade given the low cost of robotic telescopes. The very near future will see terabyte data sets of light curves routinely available to astronomers. Such data sets will contain complete samples of variable stars of all types, as well as the very rare objects or events that primarily motivate the search. It is an urgent and important problem to automate the classification of light curves in massive variability surveys.

This paper argues that new analysis methods based on neural networks will enable us to pinpoint and identify scarce transient objects in such huge data sets. Our illustrative example is the identification of scarce microlensing events against the background of variable stars. However, we envisage that the applicability of the technique is much wider.

2 MICROLENSING LIGHT CURVES

At any instant, the probability that a source star in the Galaxy shows the microlensing effect is $\lesssim 10^{-6}$. Microlensing events are hugely outnumbered by stellar variability, which is at least 100 000 times more common. The light-curve classification problem is to devise algorithms that diagnose the different kinds of variability. For applications to microlensing, the algorithm must distinguish microlensing from other sources of variability (whether intrinsic or extrinsic).

Let us assume a single, point-like, dark lens. The microlensing light curve has a characteristic form written down by Paczyński (1986). The light curve is symmetric and achromatic. As the probability of microlensing is so low, the variability must not repeat. Microlensing is readily distinguished from some, but unfortunately not all, forms of stellar variability. A cautionary history is provided by the fate of the candidate event EROS-LMC-2. This was one of the microlensing candidates uncovered by the photographic plate search of the first phase of the EROS experiment towards the Large Magellanic Cloud (Ansari et al. 1996). Although the source star of EROS-LMC-2 was known to be variable at a low level (Ansari et al. 1995), none the less microlensing seemed favoured by the excellent fit of the light curve to the data points. However, there was a substantial second bump in the light curve 8 years after the first, and EROS-LMC-2 was then discarded as a microlensing candidate (Lasserre et al. 2000).

The background in microlensing data bases is composed of periodic variables (e.g. Cepheids, RR Lyrae), eruptive variables (e.g. dwarf novae, classical novae), semiregular variables (e.g. bumpers) and the supernovae occurring in galaxies behind the source population. Of these, the most troublesome in microlensing surveys towards the Magellanic clouds and Andromeda are the bumpers and the novae-like objects. Although SNe Ia have reasonably well-understood light curves, the same is not true of other types of supernovae that can mimic microlensing rather well (for example, events 22 and 26 of Alcock et al. 2000a). Long-period bumpers may be

present as single bumps even in five seasons worth of data and they can be well fitted by the standard microlensing light curve.

Let us stress that the identification of microlensing events remains an awkward – and not fully solved – problem. For example, it probably lies at the heart of the seeming discord between the results of the MACHO and EROS experiments towards the Large Magellanic Cloud (LMC). The MACHO group identified between 13 and 17 events towards the LMC, whereas the competing EROS group found only 3 (Alcock et al. 2000a; Lasserre et al. 2000). Although the exposure times and field locations between the two experiments do vary, none the less the rate found by MACHO is at least twice that found by EROS. This same disparity is also seen in the experiments towards the Galactic Bulge, as MACHO find an optical depth to microlensing of $\sim 3.23 \times 10^{-6}$ (Alcock et al. 2000b), whereas the EROS value is about half of this (Afonso et al. 2003). Possible explanations are that the MACHO selection algorithm may be too loose (causing contamination with other variable sources) or that the EROS selection algorithm may be too harsh (causing genuine events to be discarded). It is here that neural networks may be able to make a decisive contribution.

3 AN INFORMAL INTRODUCTION TO NEURAL NETWORKS

Neural networks have been used before for pattern recognition tasks in physics (e.g. Bishop 1995). In particular, they are often used in high-energy physics experiments as triggers to select interesting events from large data sets (Müller, Reinhardt & Strickland 1995, Chapter 8). Recent astronomical applications include the classification of optical stellar spectra (Bailer-Jones et al. 1997) and galaxy type (Lahav et al. 1995), object detection in wide field imaging (Andreon et al. 2000) and predictions of astronomical time series (e.g. Conway 1998; Perdan & Serre 1998). There has also been a recent report of preliminary results on automatic light-curve classification by the *ROTSE* collaboration (Wozniak et al. 2001). An interesting review of a number of astronomical applications is given in Storrie-Lombardi & Lahav (1994).

In a neural network, the neurons are arranged in layers. The input data are fed to the bottommost layer. The output value emerges from the topmost layer, the intervening layers are hidden. The values of the neurons in any layer a_j are calculated via

$$a_j = \sum_i w_{ji} z_i, \quad (1)$$

where w_{ji} are the synaptic weights of the j th neuron with respect to the i th neuron and z_i are the activation values. The activation value is computed from the value on the neuron via an activation function g ,

$$z_i = g(a_i). \quad (2)$$

As an activation function, we use the logistic function

$$g(a) = \frac{1}{1 + \exp(-a)}, \quad (3)$$

which allows us to interpret the outputs of the network as a posteriori probabilities (Bishop 1995, Chapters 3 and 6).

We start with a sequence of input units (the ‘patterns’) for which the desired values of the output (the ‘targets’) are known. This is called the training set. Given the patterns and a set of weights, we can construct an error function E that quantifies the performance of the network. We want to obtain the weights w_{ji} that minimize the error function over the training set using a steepest-descent scheme.

We begin with random values for the weights and perform a sequence of iterative updates using a variant of back-propagation as the learning algorithm. The error derivatives with respect to the weights are

$$\frac{\partial E^n}{\partial w_{ji}} = \delta_j^n z_i^n, \quad \delta_j^n \equiv \frac{\partial E^n}{\partial a_j}, \quad (4)$$

where n labels the pattern. Using the chain rule, we obtain the back-propagation formula

$$\delta_j^n = g'(a_j) \sum_k w_{kj} \delta_k^n, \quad (5)$$

which shows how the values of δ_j^n propagate through the network, given the target value. In each iteration, the weights are updated according to the following rule:

$$\Delta w_{ij} = -\eta \sum_n \delta_j^n z_i^n, \quad (6)$$

where η is the constant learning rate. The sum is performed over all the patterns. This is equivalent to the steepest-descent method of minimizing the error. In practice, we use a refinement of this algorithm, called resilient back-propagation, which helps to prevent entrapment in local minima (see, e.g., Bishop 1995, Section 7.5.3).

As the network is converging to a minimum, it is important to prevent overtraining. This is done by feeding a different set of patterns (the ‘validation set’) to the network. The errors over the patterns in the training and the validation sets are computed separately. The training process is stopped just before the error in the validation set begins to rise. Finally, the performance of the fully trained network can be assessed with a third set of patterns (the ‘test set’). It is important to ensure that the training, validation and test sets do not contain any identical patterns.

4 IMPLEMENTATION

The experiments described below use the Stuttgart Neural Network Simulator (<http://www-ra.informatik.uni-tuebingen.de/SNNS>). Our network is composed of one input layer, one hidden layer and one output layer. The hidden layer is fully connected to the input and output layers. There are five neurons in the input layer, five neurons

in the hidden layer and one neuron in the output layer. The value of the output neuron gives the probability that the event is microlensing. The reason for the choice of five input neurons will become obvious shortly.

4.1 The training and the validation sets

There are three types of light curves in the training set: simulated microlensing events, variable starlight curves from archival sources and sample light curves from a microlensing experiment (in this case, the MACHO experiment).

Simulated microlensing events are generated by randomly choosing an impact parameter, an Einstein crossing time between 7 and 365 d and a time when the event reaches maximum. Random Gaussian noise is added with a dispersion in the range from 0.1 to 20 per cent of the maximum flux. The light curves are sparsely sampled using the MACHO sampling.

Variable stars may be divided into periodic variables and eruptive/cataclysmic variables. The former are usually easier to distinguish from microlensing than the latter, always provided more than one period can be detected in the sampled data stream. Examples of typical light curves for different types of variability are shown in Fig. 1. The periodic variables include pulsating stars (such as Cepheids and Miras) and eclipsing stars. Eruptive variables include T Tauri, S Doradus and pre-main-sequence stars. Cataclysmic variables include novae, supernovae and symbiotic variables. The relative frequencies with which these stars occur are not important in our analysis. All that matters is that the gamut of shapes is well represented in the training set. We are therefore interested as much in regular representatives as in extreme examples of the light curves. Light curves for the variable stars are selected from the sources listed in Table 1. For long data sequences, the experimental window is placed randomly on the light curve. In this way, we ensure that the bumps in the light curves do not occur in a privileged place.

Finally, there are light curves randomly chosen from the MACHO data base (specifically, from field 113 towards the Bulge). The rationale for this is that instrumental artefacts are certainly present in the MACHO light curves and it is important for the neural network to be able to recognize these.




Figure 1. This shows sample light curves of different types of variability included in the training and validation sets.

Table 1. Sources of light curves of variable stars. AAVSO is the American Association of Variable Star Observers.

Variable	Reference
Eruptive	van Genderen (1995), AAVSO
Pulsating	Antonello & Morelli (1996), AAVSO
Cataclysmic	Hamuy et al. (1996), AAVSO
Eclipsing	Brancewitz & Dworak (1980)

The training set contains 400 microlensing light curves, 150 stellar variable light curves and 200 MACHO light curves. The validation set contains the same number of light curves, although the individual representatives are obviously different. The test set are the ~ 5000 light curves from MACHO tile 113.18292, which is part of field 113 towards the Galactic bulge. Let us note that – compared with real data from a variability survey – microlensing events are overrepresented in our training and validation sets. The consequence of this is that the network will provide more false positives (as the prior probability of microlensing is too high). This is highly desirable, as the best approach to detecting such an intrinsically rare phenomenon as microlensing is to force fewer false negatives at the expense of more false positives.

4.2 Pre-processing

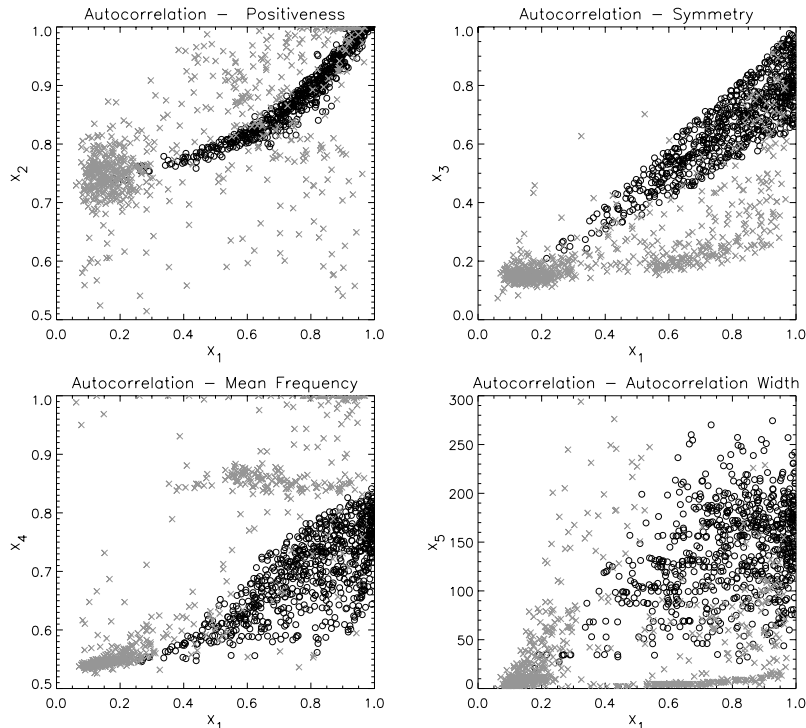
In many applications, it is both customary and advantageous to pre-process data for feeding to the neural network. The main problem with using raw photometry data is the curse of dimensionality (see Bishop 1995, Chapter 8). The simplest way of overcoming this is to extract features of the light curve and to use this as input to the network. Properly implemented, this can lead to a very efficient network, as prior knowledge can be incorporated and redundant variables can be discarded in the pre-processing. However, there

are dangers as well, as important features in the light curves can be erased.

The aim of a neural network is not to model the patterns but to model the decision boundary between the patterns. In microlensing surveys, event identification normally proceeds by making sequences of cuts, in which case the decision boundary is formed by a set of hyperplanes. The advantage of a neural network over conventional sequences of straight line cuts is that the former offers a better chance of describing a complicated decision boundary accurately.

Microlensing events are characterized by the presence of an (i) excursion from the baseline, that is (ii) positive, (iii) symmetric and (iv) single. The event itself is characterized by (v) a time-scale. Motivated by these five features, we extract from the light curves the following five parameters, which are inputs to the neural networks.

The first parameter, x_1 , is the maximum value of the autocorrelation function. This helps to discriminate against noise and to identify the presence of any signal. The second parameter, x_2 , is calculated as follows. First, we compute the median of the flux measurements, which gives a good approximation to the baseline. We then compute the mean of the data points lying above and below the median and finally take their ratio. This is then mapped on the interval $[0.5, 1]$ with the logistic function. The input x_2 tests for the positiveness of the excursion. The third parameter, x_3 , is the maximum value of the cross-correlation function of the light curve with the time-reversed light curve. This provides a test for symmetric events. The fourth parameter, x_4 , is the mean frequency $\langle \nu \rangle$ calculated using the power spectrum $P(\nu)$ as a weighting function. For a periodic variable, we expect a shift in the weighted mean frequency from zero. We compress x_4 using the logistic function to lie in the range $[0.5, 1.0]$. Finally, the fifth parameter, x_5 , is the width of the autocorrelation function, as judged by its standard deviation. If the event is microlensing, then the width is a rough indication of the time-scale.

**Figure 2.** This shows projections on to the principal planes of the five-dimensional space of inputs. Bold circles show the microlensing events and grey crosses the variable stars and noise in the training and validation sets.

To motivate this choice of inputs, Fig. 2 shows the locations of all the patterns in the validation and training sets. The desideratum is that the choice of inputs offers a clear separation between microlensing events and other patterns in the five-dimensional space (x_1, \dots, x_5). The projections of this space on to the principal planes offer grounds for believing this, as there is already good partial separation in some of the plots (e.g. x_1 versus x_3) and good evidence for regularities in others (e.g. x_1 versus x_2). The final proof that the choice of inputs is good can, however, only be provided by the performance of the network on the test set.

Note that Fig. 2 plots unnormalized input variables; however, the neural network uses normalized inputs. Scaling of the inputs to numbers of the order of unity is often useful, as this means that the network weights also typically take values of the same order (Bishop 1995, Chapter 8). Pictorially, this can be thought of as requiring the hyperplanes associated with each hidden unit to intersect close to the origin and near the centre of the datacloud. For each input variable, this scaling is done by subtracting the mean and dividing by the standard deviation to give the normalized inputs.

So far, we have skirted round the problem of missing data points. For MACHO data, ~ 10 per cent of the light curves have gaps of the order of a few days (aside from the 5-month gaps when the Galactic bulge is not visible from Australia). To compute the correlation functions, the data are treated as if they were uniformly sampled. This gives rise to some errors. If the typical gap size is much smaller than the event time-scale, then any errors we have introduced by this procedure will be small. If the gap size relative to the time-scale is very large, then no classification can be plausibly extracted. If the gap size is of the same order as the time-scale, then the experiment needs redesigning. The input most sensitive to missing data is x_4 because this requires computation of the power spectrum. There are, however, existing algorithms to do this for unevenly sampled data (e.g. Lomb's periodogram as implemented by Press & Rybicki 1989; Press et al. 1992), which we employ.

Note that pre-processing gives rise to fast and powerful neural networks, but it can also cause loss of potentially important information in the data. To check this, we can allow a neural network itself to perform the projection. This leads to much bigger neural networks, which consequently take longer to converge. However,

it does have the advantage that no assumptions are built in from the beginning. In this spirit, we experimented with a big neural network, which takes as the two input layers the unadulterated flux measurements and errors at the sampling times and has ~ 200 hidden neurons. Once converged, the performance of this big network is similar to the performance of smaller networks on pre-processed data. From this, we draw the conclusion that our pre-processing has not caused any serious degradation of the information in the data.

4.3 Training

In training, the weights are initialized to random values. We then perform iterations to reduce the error function

$$E^n = - \sum_n [t^n \log y^n + (1 - t^n) \log(1 - y^n)], \quad (7)$$

where t^n and y^n are the target and the response of the output neuron for the n th pattern. We have chosen this form of the error function (the so-called *cross-entropy* error function) as appropriate for two class problems (see, e.g., Bishop 1995, Section 6.7). Given our choice of activation (3) and error functions (7), the output y^n approximates the posterior probability $P(\text{microlensing} | \text{inputs})$.

The neural network must be able to generalize from the patterns in the training set, and not merely reproduce them. A worry is that the network will be overtrained and will reproduce structures of the decision boundary in unnecessary detail. To guard against this, we use early stopping as illustrated in Fig. 3. The performance of the network on the validation set is compared with that on the training set and the training stopped just before the error in the validation set rises. Another safeguard is provided by the introduction of a small amount of noise to the weights on each iteration, which guards against entrapment in a local minimum.

If training is started from different initial weights, we converge to slightly different final weights. This makes it advantageous to use a committee of 10 neural networks (Bishop 1995, Section 6.7). There are a total of 1500 light curves available. For each member of the committee, the 1500 patterns are split in half randomly to give validation and training sets with 750 members. For each pattern, the final output is the average of the output of all 10 neural networks.




Figure 3. This shows the value of the cross-entropy error function versus the epochs of training (number of iterations) for the patterns in the training and validation sets. The long-dashed line shows the point at which the training is stopped. The sum of errors is ~ 15 out of the 700 patterns in the set.




Figure 4. This shows the histogram of output values for the 1500 patterns in the validation and training sets. Note the clean separation between microlensing and other types of variability.

The histogram of the output values for the combined validation and training set is shown in Fig. 4. There is a very clean separation of microlensing events and other forms of variability. The non-microlensing events are strongly peaked at a probability $y = 0$, but there are a few events (~ 10) that extend up to $y = 0.2$. The microlensing events are strongly peaked at $y = 1$, although again there are a few (~ 10) that extend down to $y = 0.7$. The probability $y = 0.5$ corresponds to the formal decision boundary (Bishop 1995, Section 10.3). In fact, between $0.2 < y < 0.7$, there are almost no events in the histogram. If, when presented with a light curve, the neural network does give an output in this range, then the classification is in reality uncertain. This is because any error in the output can cause it to straddle the formal decision boundary. This range of outputs really corresponds to patterns that are not present in the training and validation sets. This is valuable as it offers the possibility of the detection of unexpected and novel events in variability surveys.

There are just three microlensing events out of 800 that are misclassified (i.e. have $y < 0.5$). These are scarcely visible on the histogram. It is interesting to locate these events in our input space (see Fig. 2). These events have input coordinates (1.05, 0.74, 0.16, 0.54, 20.63), (0.19, 0.73, 0.16, 0.55, 34.53) and (0.22, 0.74, 0.21, 0.55, 34.53). They are small-amplitude or short-duration events dominated by noise, as indicated by the value of the x_1 input, which measures the presence of the signal. There is one false positive (i.e. a non-microlensing light curve with $y > 0.5$), which has coordinates (0.27, 0.75, 0.21, 0.56, 40.87). This is a light curve from the MACHO tile, which is probably neither microlensing nor a variable star, but just noise.

5 TESTS TOWARDS THE BULGE FIELDS

5.1 Normal events

All MACHO light curves extracted with conventional point spread function (PSF) photometry (such as SODOPHOT) are now publically available (Allsman & Axelrod 2001). As a first test, we use light curves from tile 18 292 of field number 113, which lies towards the

Galactic bulge. This tile contains ~ 5000 light curves of which one was identified by MACHO as a microlensing event. The MACHO data are taken at a site with moderate seeing. According to Alcock et al. (2000b), the median seeing is ≈ 2.1 arcsec. This means that the quality of the data is sometimes quite poor. To allow for this, we clean the light curves by removing all isolated points with more than 3σ deviation from the immediately preceding and succeeding data points. In general, this makes good sense as it removes outliers, but it can sometimes remove meaningful data points for very rapid brightness variations.

Each cleaned light curve is shown to the committee of neural networks. The red and blue passband data are analysed separately. In principle, it would be advantageous to analyse the red and blue data together because most variable stars show chromaticity differences. However, this option is not open to us at the moment because the publically available colour information on variable stars is still quite limited. Fig. 5 shows the results of the deliberations of the committee. The probability of microlensing given the blue data is shown against the probability given the red data. There is only one pattern that satisfies $y > 0.7$ in both colours, namely the event identified by MACHO as BLG-95-1. It is clearly and cleanly separated from the rest of the patterns in the figure as a black circle in the topmost right-hand corner. There is an additional pattern that has output values of $y \approx 0.6$ for both the red and blue data. This falls within the regime of novelty detection. Its input coordinates are (0.43, 0.70, 0.26, 0.61, 183.14). It is a very long event since its x_5 value is higher than typical values for microlensing. It falls into a poorly sampled region in Fig. 2, which suggests why this low signal-to-noise ratio light curve was dragged into the microlensing range. Its light curve is shown in the upper panel of Fig. 6. It is most probably a form of stellar variability that does not lie in the training and validation sets. It is interesting to note that there are a number of light curves with output greater than 0.9 in one band, but not in the other. Shown in the lower panel of Fig. 6 is a typical example, in this case securely identified in blue ($y > 0.95$) but not in red ($y < 0.05$). The blue light curve does indeed look like a microlensing event, but the better sampling in the red passband shows a highly active many-humped light curve, which is most probably an eruptive variable.




Figure 5. This shows the output of the committee of neural networks for all light curves in tile 113.18292, which is publically available from the MACHO project website (see Allsman & Axelrod (2001)). Shown on the vertical and horizontal axes are the probabilities that the blue and red light curves are microlensing. There are ~ 5000 light curves on the tile, including one event BLG-95-1 identified by the MACHO collaboration as microlensing. This is shown as the black spot.

As a second test, we analyse the light curves for all 36 events in Alcock et al. (2000b) that were identified on the basis of conventional PSF photometry. Table 2 shows the results of the poll of the committee. In each case, the output of the neural network on the red and the blue data is given. Of course, it is important to bear in mind that the classification algorithm used by the MACHO group is itself probably not 100 per cent efficient. There are reasons to believe – both from the very high rate towards the Galactic Centre, which is incompatible with theoretical models of the Galaxy, and from the differences between the MACHO and EROS results – that the subsample of candidates found by MACHO may have some contamination. There are a total of 19 events identified with a probability $\gtrsim 0.5$ as microlensing in both the red and blue filters. In fact, these events are all beyond reproach as microlensing candidates as the probability is $\gtrsim 0.9$. Events 28 and 36 are securely identified in the blue data, but the red data are corrupted. Events 4, 9, 15, 20, 22 and 25 are identified in the red data, but not in the blue. Lastly, there are nine events for which no microlensing signal whatsoever is detected (event numbers 3, 5, 13, 17, 21, 24, 29, 31, 32). We shall examine the light curves of some of these events shortly, but for the moment let us emphasize that there is no guarantee that the original identification by the MACHO collaboration was correct.

Fig. 7 shows the contours of probability for the training and validation sets in the input space. Light grey means that the probability is greater than 0.5 and corresponds to the formal decision boundary (see Bishop 1995 Section 10.3). Dark grey means that the probability is greater than 0.9 and corresponds to almost certain microlensing. The irregularity of the contours is a result of the fact that some regions are poorly sampled in the training and validation sets. The contours have been drawn with a view to guiding the eye. Super-

posed on the contours in Fig. 7 are the events. The nine unfilled circles are those identified by the network as variable stars but by MACHO as microlensing events. The black circles are those for which both MACHO and the network agree as microlensing.

There are a number of things to note in the diagram. First, it is evident that the network has the ability to extrapolate from the validation and training sets and assign relative importance to the combinations of features extracted by the input variables. This is clear because there are events securely identified although they lie outside the contours (for example, event 18 is unambiguously identified despite lying outside the probability contours in the two top panels). Secondly, the x_4 input is the only one for which explicit allowance has been made for noise and sampling. The network seems to assign greater importance to this input, as almost all the filled circles lie within the projected 90 per cent probability contour. This suggests that further improvements may be possible by allowing for noise in the extraction of other input parameters (for example, using extrapolation for the correlation analysis). Thirdly, the separation between the 0.5 and 0.9 probability contours is typically very small, so the contour surface is very steeply rising. Such outputs can correspond to novelty detection. Accordingly, they occupy only a small region of the input space and so novelty detection occurs – as is highly desirable – for only a few light curves. The small separation between the contours provides justification for the sizes of the training and validation sets. If there are too few patterns in these sets, then the separation would widen. Such widening happens in our network only in a few unimportant regions, which are physically inaccessible (that is, such a combination of input variables gives rise to light curves that do not occur in nature). Fourthly, all the unfilled circles have $x_1 < 0.4$ and so lie in the noise-dominated regime. However, the values of x_2 indicate the presence of substantial positive excursions. This is already enough to tell us that the noise in the MACHO data is strongly non-Gaussian.

Fig. 8 shows the light curves for eight of the events corresponding to the unfilled circles. For some of these events, it looks as though there are secondary bumps (e.g. event 3). For others, the bump is not properly contained (e.g. events 5 and 31) or the bump is overwhelmed by noisy data (e.g. events 13 and 17). It seems that the performance of our network is excellent, as these events certainly need to be looked at with care before accepting a classification as microlensing. However, it is premature to conclude that MACHO have misclassified these events. This is because the MACHO group have reprocessed all the light curves with difference image analysis (DIA) and this will improve the quality of the light curves, reducing noise and contamination from nearby stars. However, without having the DIA light curves, we cannot confirm their verdict of microlensing.

5.2 Exotic events

Some microlensing light curves can show deviations from the standard Paczyński form caused by parallactic or finite-source size effects or by binarity and so on (see, e.g., Mao & Paczyński 1991; Mao & di Stefano 1995; Kerins & Evans 1999; Mao et al. 2002). In Table 3, all the exotic events identified in Alcock et al. (2000b) using the SODOPHOT photometry package are processed with the committee of neural networks.

Parallactic events (such as 96-BLG-12) occur when the Einstein radius projected on to the observer's plane is of the order of an astronomical unit. In such a circumstance, the changing motion of the Earth around the Sun during the event is detectable by an asymmetry in the light curve with respect to the peak. Events showing deviations




Figure 6. The upper panels show the blue and red light curves for the event identified by the grey spot in Fig. 5. The lower panels show the light curves for the event securely identified in the blue band (left-hand panel), but not in the red band (right-hand panel). In all cases, the horizontal axis is time in days, the vertical axis is flux in ADU s^{-1} . The vertical dashed lines mark 5-month gaps when the bulge is not visible from Australia.

Table 2. This shows the output of the committee of neural networks on the subset of candidates towards the bulge in Alcock et al. (2000b) which are selected on the basis of the conventional PSF photometry package (SODOPHOT). The results of the analysis of the red and blue light curves are shown separately. The output is the probability that the event is microlensing. (Note that the red data for event 17 and the blue data for event 27 are unavailable.)

ID	MACHO ID	R	B	ID	MACHO ID	R	B
1	97-BLG-24	0.93	1.00	2	95-BLG-5	0.92	0.96
3	97-BLG-42	0.00	0.00	4	97-BLG-s4	0.72	0.00
5	95-BLG-15	0.00	0.00	6	95-BLG-s8	1.00	1.00
7	97-BLG-18	1.00	1.00	8	96-BLG-26	1.00	1.00
9	97-BLG-38	0.64	0.47	10	97-BLG-58	1.00	1.00
11	96-BLG-1	1.00	1.00	12	97-BLG-2	1.00	1.00
13	96-BLG-14	0.00	0.00	14	95-BLG-s9	0.99	0.91
15	96-BLG-21	0.68	0.00	16	95-BLG-1	1.00	1.00
17	96-BLG-s10	—	0.16	18	96-BLG-20	0.99	1.00
19	96-BLG-10	0.99	0.90	20	95-BLG-4	0.81	0.02
21	95-BLG-23	0.00	0.00	22	95-BLG-s13	0.64	0.05
23	95-BLG-10	1.00	1.00	24	97-BLG-4	0.00	0.00
25	97-BLG-16	1.00	0.23	26	96-BLG-8	0.96	0.99
27	95-OGLE-16	0.99	—	28	95-BLG-39	0.16	1.00
29	95-BLG-3	0.39	0.00	30	97-BLG-37	1.00	1.00
31	97-BLG-14	0.21	0.00	32	95-BLG-11	0.01	0.00
33	96-BLG-31	0.81	1.00	34	96-BLG-s16	1.00	0.83
35	97-BLG-s14	0.80	0.90	36	95-BLG-22	0.30	0.79

caused by finite source size (such as 95-BLG-30) occur whenever the angular size of the source is of the same order of magnitude as the angular Einstein radius. They are usually flatter-topped than the classical Paczyński curves for microlensing by a point source.

For both of these kinds of deviation, the neural network committee performs well, as shown in Table 3. All the parallaxic and finite source size events are identified as microlensing.

However, binarity can cause much substantial deviation. For example, strong binary events have additional peaks, although these can sometimes be missed if sampled irregularly. Weak binary events may just have distortions to the peak or the wings of the light curve. Accordingly, we might expect the detection of binary light curves to require the training and testing of a new neural network. This is supported by the results in Table 3. Here, 96-BLG-12 is identified by the committee, whereas 96-BLG-3 falls into the domain of novelty detection. It is reassuring that in the former case, the event is recognized, while in the latter case, the event is recognized as a new phenomenon. The development of software to recognize binary events is a problem that has not been fully solved by any of the microlensing collaborations to date. It seems reasonable to expect neural networks to play a powerful role here.

6 CONCLUSIONS

This paper has devised a working neural network that can distinguish simple microlensing light curves from other forms of variability, such as eruptive, pulsating, cataclysmic and eclipsing variables. The network is structured to have five input neurons and one output neuron. The inputs and output are separated by a layer of hidden neurons. The simplicity of the network means that it can be trained very quickly and it can be used to process huge data sets in less than a second. Each light curve is pre-processed to provide five inputs to be fed to the network. In our application, the five inputs were chosen on physical grounds as good discriminants for microlensing. In other

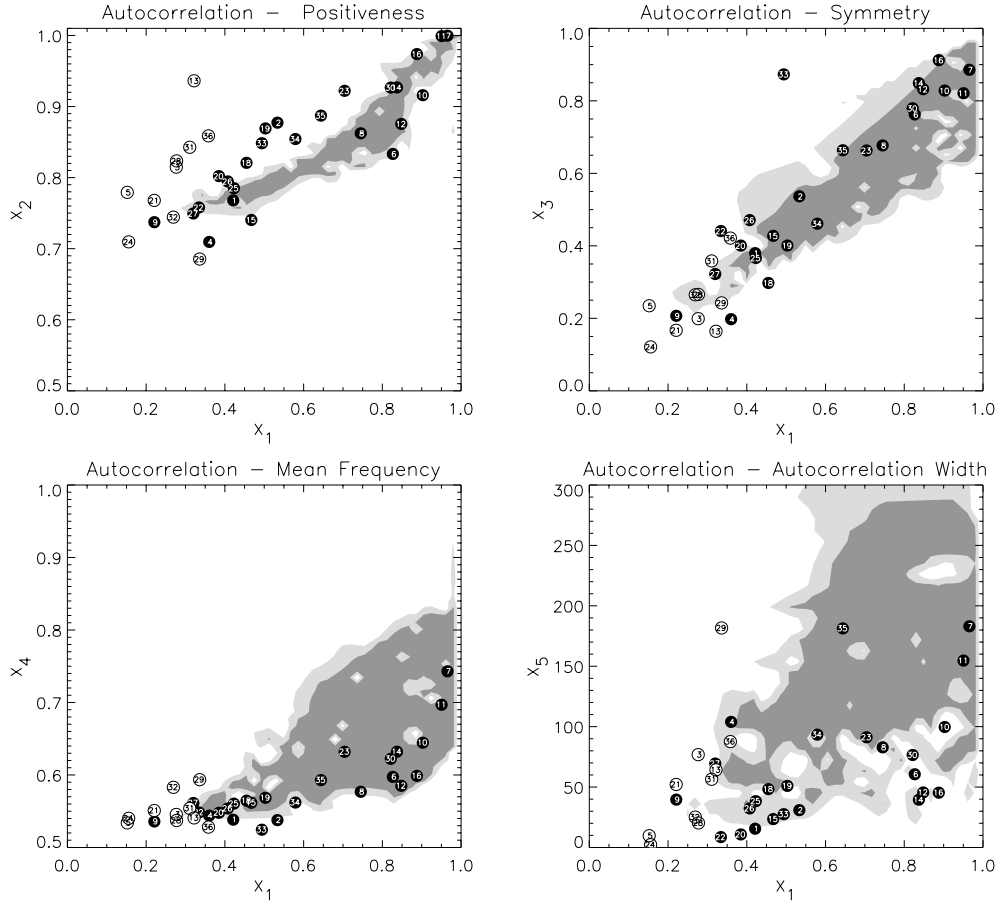


Figure 7. The grey-scale contours show the probability of microlensing in the input space (x_1, \dots, x_5) as judged from the patterns in the training and validation sets. The circles show the locations of the microlensing events identified by MACHO using conventional PSF photometry. Filled circles designate the events also identified in the red filter by the network. Unfilled circles are not identified. Numbers within circles refer to our event designations in Table 2. (Light grey means that the probability is greater than 0.5 and dark grey greater than 0.9.)

applications, different input variables may be optimum. Our network has been constructed so that the output is the posterior probability of microlensing.

We believe that neural networks offer three important advantages over conventional techniques using in microlensing experiments. First, the decision boundary separating microlensing from non-microlensing may be rather complicated. At present, all microlensing collaborations use a series of cuts (for example, on the goodness of fit to a Paczyński curve, on achromaticity and so on). This is the crudest form of decision boundary. However, even simple neural networks can reproduce complicated decision boundaries and so the technique is both more efficient and more flexible. Moreover, once a light curve has failed to pass a cut at the early stages of a conventional selection process, it is lost for any further analysis. However, neural networks assign a relative importance to the input parameters, thus the decision is based on the whole of the information available.

Secondly, neural networks offer a superior way of calculating the event rate, avoiding the need for any kind of efficiency calculation. The classical procedure of identifying events with cuts is inefficient, and this necessitates the cumbersome Monte Carlo calculation of the numbers of synthetic events passing the cuts. However, a properly designed neural network can reproduce the decision boundary well and can enable the event rate to be computed directly for comparison

with theoretical models, thus completely sidestepping the need for any Monte Carlo calculation of the efficiencies.

Thirdly, novelty detection is made both more precise and easier by neural networks. The conventional approach relies on an examination by eye of the events left over after applying a sequence of cuts. For our neural network, we have argued that all light curves with outputs between 0.2 and 0.7 may be examples of light curves not contained within the training set. These are the events that need looking at very carefully. With the even more massive data sets of the future, it will be important to identify possible novel events as quickly and as efficiently as possible.

From the point of view of microlensing, it is interesting to extend the work in this paper to include additional effects. Some of the ongoing microlensing experiments are working in the highly blended regime. For example, the POINT-AGAPE (Paulin-Henriksson et al. 2002, 2003), WeCAPP (Riffeser et al. 2001) and MEGA (Crotts et al. 2000) collaborations are all monitoring the nearby galaxy M31. Here, the individual stars are not resolved, so the flux in a pixel or superpixel is followed (Baillon et al. 1993). The range of light curves in such pixel lensing experiments is very wide – for example, microlensing events can occur in the same superpixel as bright variable stars (e.g. the event PA-99-N1 described in Aurière et al. 2001). So, the identification of microlensing events becomes still more daunting. As the complexity of the pattern recognition

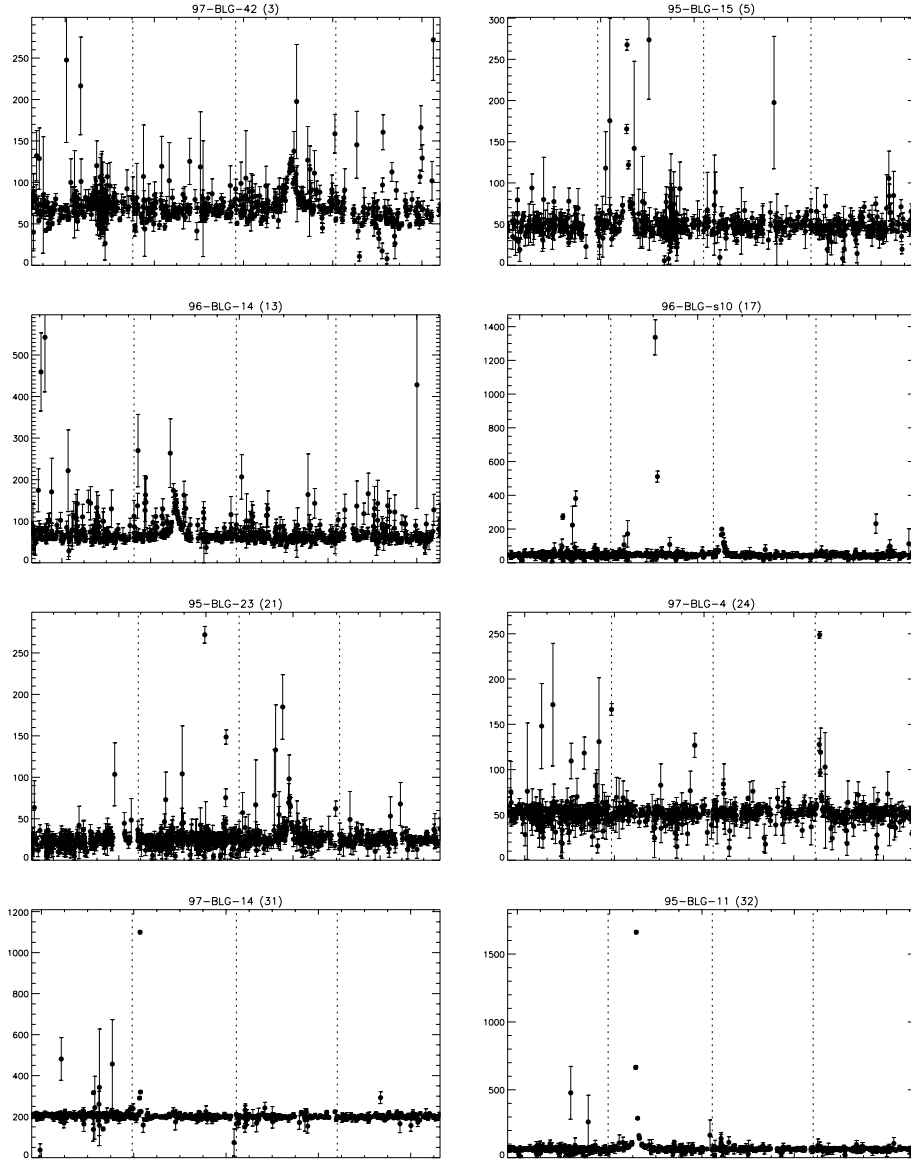


Figure 8. This shows the data points for eight of the events classified as non-microlensing by the network and as microlensing by the MACHO collaboration. The vertical axis is flux in ADU s^{-1} and the horizontal axis is time in days. The data is presented as four strips of 7-month sequences; the 5 months when the bulge is not visible from Australia is marked by the vertical dashed lines.

task increases, so we expect the power and flexibility of the neural network approach to pay increasing dividends. Also, in this paper, we have concentrated on the microlensing data sets towards the bulge, for which the source stars are often bright. It is important to apply our techniques to the microlensing events towards the Large Magellanic Cloud. Here, the task is harder as the source stars are fainter and there is serious contamination from supernovae in background galaxies. This work will be the subject of a separate publication.

Although our application has been strongly focused on microlensing, the technique is of general applicability in astronomy. There are numerous ongoing or planned massive photometry surveys using robotic telescopes (*ROTSE*), wide field cameras (*WASP* and *VISTA*) and space-borne satellites (*GAIA* and *Eddington*). Although the goal of the surveys is different, the basic method is the same – a brute-force search through many terabytes of data for interesting but rare

events, whether they are planetary transits, cataclysmic variables or optical flashes. We envisage such tasks being routinely devolved to neural networks in the astronomy of the future. In each case, cascades of neural networks could be trained to filter and identify the various classes of variable stars, to pinpoint the target events of interest and to isolate the unexpected or new classes of phenomenon, which need looking at very carefully.

7 SPECULATIONS

Suppose the goal is to monitor the whole sky for variability at short time intervals down to 20th magnitude (roughly a billion objects in our Galaxy). In this speculative final section, we ask what is possible now and what will be possible by 2010?

Let us consider the simple situation of a single neural network program running on a single computer. The middle-range

Table 3. This shows the output of the committee of neural networks on the exotic events identified towards the bulge in Alcock et al. (2000b). These are all exotic events selected on the basis of the conventional PSF photometry package (SODOPHOT); f denotes deviations caused by finite source size, p owing to parallactic effects and b owing to binarity. The output is the probability that the red and blue data correspond to a microlensing event.

ID	MACHO ID	Deviation	<i>R</i>	<i>B</i>
1	95-BLG-30	f	1.00	1.00
2	96-BLG-12	p	1.00	0.99
3	97-BLG-1	b	0.95	1.00
4	97-BLG-8	p	1.00	1.00
5	97-BLG-26	p	1.00	1.00
6	96-BLG-3	b	0.83	0.51
7	95-BLG-18	p	0.99	0.75

hardware situation today is typically a processor running at 2200 MHz [corresponding to approximately 1000 million instructions per second (MIPS)]. In order to predict the situation in 2010, we can use ‘Moore’s law’, which says that the numbers of transistors in a processor chip doubles every year or so. Thus, by 2010 the processor speed should be $\sim 100\,000$ MIPS, compared with about 1000 MIPS today. However, to evaluate the progress in run time, we must consider both hardware and the compiler. Benchmarking programs such as SPEC provide us with some clues as to what will be achieved in 2010. If we look at the evolution in performance results on SPEC tests for computers between 1995 and 2000, we find an approximate speed-up factor of 16, or roughly 1.74 per year. We can extrapolate this progress over the 2002–2010 period, which gives a speed-up factor of 85. In other words, both Moore’s law and the extrapolation of benchmarking suggest rather similar speed-up factors of roughly two orders of magnitude by 2010.

The time required to run the neural network itself is negligible compared with the time required to run the pre-processing, which extracts the parameters used by the neural network. Our present pre-processing program requires 10^{-4} s to analyse 100 data points for a single star. We have chosen 100 data points as it might correspond to sampling three times a night for 1 month, which is reasonable for the detection of fast transient events. Alternatively, it might correspond to sampling once a night for 3 months, which is reasonable for the detection of variability such as microlensing with a characteristic time-scale of ~ 1 month. At present, it therefore takes $\sim 10^5$ s (or over a day) to analyse such a data set for the whole sky. By 2010, it will take only 20 min for such a program to run on the whole sky (using the speed-up factor of 85). More generally, the time taken in seconds to analyse a set of N_{pts} data points for N_* stars in 2010 is

$$t \sim 2 \times 10^{-9} N_* N_{\text{pts}} \log N_{\text{pts}}. \quad (8)$$

Let us assume there are 8 h of observing time a night and that we wish to process a month of data for the whole sky in real time. Then we can derive the real-time equation

$$t^2 \sim 1.7 \times 10^6 (13.7 - \log t), \quad (9)$$

which has a solution $t \sim 50$ min. In other words, real-time processing of variable phenomena across the entire sky down to 20th magnitude will be possible for sampling rates of $\gtrsim 1$ h by 2010.

Our speculative calculation errs on the pessimistic side because we have not taken into account any correction for application of parallel processing or the rapidly developing GRID technology for

high-performance computing. However, it surely does enough to convince the reader that, properly trained, neural networks can analyse huge data sets very quickly. This will become one of the methods of choice for data-mining in the massive variability surveys of the very near future.

ACKNOWLEDGMENTS

VB is supported by a Dulverton Scholarship. YLD thanks PPARC, while NWE thanks the Royal Society for financial support. In this research, we have used, and acknowledge with thanks, data from the American Association of Variable Star Observers (AAVSO) International Data base based on observations submitted to the AAVSO by variable star observers worldwide. This paper also utilizes public domain data obtained by the MACHO Project, jointly funded by the US Department of Energy through the University of California, Lawrence Livermore National Laboratory under contract no W-7405-Eng-48, by the National Science Foundation through the Centre for Particle Astrophysics of the University of California under cooperative agreement AST-8809616, and by the Mount Stromlo and Siding Spring Observatory, part of the Australian National University. In this respect, we particularly wish to thank Robyn Allsman and Tim Axelrod for help in acquiring data for an entire tile. We also thank the anonymous referee for a helpful report.

REFERENCES

- Afonso C. et al., 2003, *A&A*, in press (astro-ph/0303100)
 Alcock C. et al., 1997, *ApJ*, 486, 697
 Alcock C. et al., 2000a, *ApJ*, 542, 281
 Alcock C. et al., 2000b, *ApJ*, 541, 734
 Allsman R.A., Axelrod T.S., 2001, preprint (astro-ph/0108444)
 Andreon S., Gargiulo G., Longo G., Tagliaferri R., Capuano N., 2000, *MNRAS*, 319, 700
 Ansari R. et al., 1995, *A&A*, 299, L21
 Ansari R. et al., 1996, *A&A*, 314, 94
 Antonello E., Morelli P.L., 1996, *A&A*, 314, 541
 Aurière M. et al., 2001, *ApJ*, 553, L137
 Bailer-Jones C.A.L., Irwin M., Gilmore G., von Hippel T., 1997, *MNRAS*, 292, 157
 Baillon P., Bouquet A., Giraud-Heraud Y., Kaplan J., 1993, *A&A*, 277, 1
 Bishop C., 1995, *Neural Networks for Pattern Recognition*. Oxford Univ. Press, Oxford
 Branczewitz H.K., Dworak T.Z., 1980, *Acta Astron.*, 30, 501
 Conway A.J., 1998, *New Astron. Rev.*, 42, 343
 Crotts A., Uglesich R., Gould A., Gyuk G., Sackett P., Kuijken K., Sutherland W., Widrow L., 2000, in Menzies J.W., Sackett P.D., eds, *ASP Conf. Ser. Vol. 239, Microlensing 2000: a New Era of Microlensing Astrophysics*. Astron. Soc. Pac., San Francisco, p. 318
 Hamuy M., Phillips M.M., Suntzeff N.B., Schommer R.A., Maza J., Smith R.C., Lira P., Aviles R., 1996, *AJ*, 112, 2438
 Kerins E.J., Evans N.W., 1999, *ApJ*, 517, 734
 Lahav O. et al., 1995, *Sci*, 267, 859
 Lasserre T. et al., 2000, *A&A*, 355, L39
 Mao S., di Stefano R., 1995, *ApJ*, 440, 22
 Mao S., Paczyński B., 1991, *ApJ*, 374, L37
 Mao S. et al., 2002, *MNRAS*, 329, 349
 Müller B., Reinhardt J., Strickland M.T., 1995, *Neural Networks: an Introduction*. Springer-Verlag, Berlin
 Paczyński B., 1986, *ApJ*, 304, 1
 Paczyński B., 2001, in Banday A.J., Zaroubi S., Bartelmann M., eds, *Mining the Sky*. Springer-Verlag, Berlin, 481 (astro-ph/0110388)
 Paczyński B., 2002, in Chen W.P., Lemme C., Paczyński B., eds, *ASP Conf. Ser. Vol. 246, Small Telescope Astronomy on a Global Scale*. Astron. Soc. Pac., San Francisco, in press (astro-ph/0108112)

- Paulin-Henriksson S. et al., 2002, *ApJ*, 576, L121
 Paulin-Henriksson S. et al., 2003, in press (astro-ph/0207025)
 Perdang J., Serre T., 1998, *A&A*, 334, 976
 Press W.H., Rybicki G.B., 1989, *ApJ*, 338, 277
 Press W.H., Teukolsky S., Vetterling W.T., Flannery B., 1992, *Numerical Recipes*, 2nd edn. Cambridge Univ. Press, Cambridge, Ch. 13
 Riffeser A. et al., 2001, *A&A*, 379, 362
 Storrie-Lombardi M.C., Lahav O., 1994, *Vistas Astron.*, 38, 249
 Udalski A., Szymanski M., Kaluzny J., Kubiak M., Mateo M., Krzeminski W., Paczyński B., 1994, *Acta Astron.*, 44, 227
 van Genderen A.M., 1995, *A&A*, 304, 415
 Wozniak P.R. et al., 2001, *American Astronomical Society Meeting*, 199, 130.04

This paper has been typeset from a \TeX/L\AA\TeX file prepared by the author.